

Supplementary Material: Cross-Age Face Verification from Mined Same-Person Social-Media Pairs

Anonymous submission

This appendix reports the controls, audits, full operating-point tables, and the data/model card deferred from the main paper for space. Numbering here is internal to this document; references to the main paper are by name. All experiments use the curated data and the same weak-backbone attribution protocol described in the main paper.

1 Extended Related-Work Positioning

Table 1 positions our naturally-supervised source against the discriminative, contrastive and synthetic-aging lines of prior work. The common thread is that prior methods depend on *manual* identity and/or age labels or on synthesized cross-age samples, whereas we mine the structure “one post = one person at different ages” as a label-free positive-pair generator.

2 Data and Audit Details

Age extraction. Across all 33,697 captioned multi-photo posts, the LLM and regex agree on 72.8% (Table 2); the LLM fills a further 10.8% the regex missed and is favored in the bulk of the 13.8% mismatches (it avoids reading calendar years or year-gaps as ages).

Group integrity. Person-clustering had already separated most multi-person posts; an LLM group-integrity pass (Table 3) flags only 421 person-groups as noisy (multi-person/meme).

Dedup-threshold robustness. The -52% positive-count reduction from near-duplicate removal is not an artifact of the 0.97 cutoff: it is stable across 0.93–0.97 and degrades only at the very strict 0.99 (Table 4).

Small manual validation. One annotator spot-checked the automatic pipeline (Table 5). With a single annotator and no inter-annotator κ this *supports* but does not *certify* the supervision; a large multi-annotator audit remains beyond our resources, so residual label noise is acknowledged as a limitation. Samples are drawn from the curated corpus and were not used to tune thresholds or filtering rules.

3 Objective Comparison (Full)

Figure 1 plots the objective comparison summarized in the main paper: the contrastive, ArcFace and CosFace objectives coincide on cross-age transfer (within ± 0.006 on FG-NET large-gap), a noise-robust sub-center ArcFace adds nothing

This is an anonymized submission for review purposes only. Distribution, citation, or public sharing of this manuscript is strictly prohibited. Copyright and publication details will appear in the final version if accepted.




Figure 1: Objective comparison: contrastive, ArcFace and CosFace coincide on cross-age transfer; SphereFace trails. ArcFace/CosFace forget slightly less on easy benchmarks.

on the curated pairs, and only SphereFace—hardest to optimize on 2–3-shot identities without λ -annealing—trails. ArcFace/CosFace additionally forget slightly less on the easy benchmarks, a practical recommendation.

4 Headroom and Data-Scaling Tables

Table 6 gives the numeric headroom curve (the gain decreases monotonically with frozen strength, with a clean AdaFace IR-50→IR-101 depth control), and Table 7 the data-scaling law (the gain saturates by $\approx 10\%$ of identities).

5 Mechanism: Age-Leakage Probe and Disentanglement

A linear probe (Table 8) shows apparent age stays well decodable from the identity embedding for *all* models (≈ 0.63 – 0.65 balanced accuracy $\gg 0.25$ chance) and changes only slightly, while identity AUC rises sharply—the method learns an invariant *metric* (the cosine stops relying on the age axis), not

Work	Supervision	Mechanism	Our difference
OE-CNN (2018) (Wang et al. 2018)	labeled id+age	orthogonal feature decomposition	no manual id/age labels, not decomposition
DAL (2019) (Wang et al. 2019)	labeled id+age	adversarial id/age factorization	mined pairs + simpler objective
MTLFace (2021) (Huang, Zhang, and Shan 2021)	labeled id+age	recognition + age synthesis (multi-task)	weaker pipeline; novel supervision origin
CACon (2023) (Wang, Sanchez, and Li 2024)	semi-sup. + synthetic	contrastive on synthesized cross-age	real mined pairs > synthetic aging
Synthetic Ageing (2024) (Yao et al. 2024)	synthetic	train on aged images	we make <i>real</i> mining the thesis
OrdCon (2025) (Wang, Sanchez, and Li 2025)	age-ordered contrastive	con- order-enhanced features	mine natural pairs, no explicit age order

Table 1: Positioning vs. prior work.

Category	<i>n</i>	Share
Agree (same ages or both empty)	24,531	72.8%
LLM filled (regex missed)	3,656	10.8%
Age mismatch	4,658	13.8%
LLM missed (regex found)	852	2.5%

Table 2: Age-extraction audit: LLM vs. regex (33,697 captioned multi-photo posts).

Category	<i>n</i>	Share
Single (one person across ages)	31,275	92.8%
Multi-person (different people)	2,034	6.0%
Unknown	336	1.0%
Meme / collage	52	0.2%

Table 3: LLM group-integrity classification (33,697 posts).

an age-free representation. Explicit adversarial disentanglement (gradient reversal + age head, MTLFace-style) reduces leakage no further and adds only a small cross-age gain over +pairs (+0.006 FG-NET large-gap, +0.015 our.25+); its absolute large-gap (≈ 0.853) matches the ArcFace objective (0.851), reinforcing that the data—not the objective or the disentangling add-on—sets the cross-age ceiling.

6 Apparent-Demographic Audit (Full)

Stratifying by *apparent* gender/age of the pair anchor (estimated, not self-identified), no group degrades and every gain is significant at the 95% paired-bootstrap level (Table 9, Fig. 2). The largest gain is for the youngest group 0–17 (+0.130), where the frozen baseline is weakest. We avoid the word “fair” as a solved property: the source is apparent-female-skewed (76.7%), attributes are apparent (not ethnicity

Dedup cos	Near-dup faces	Positives	Reduction
0.93	8,653	31,304	−52.5%
0.95	8,564	31,611	−52.0%
0.97 (used)	8,304	32,360	−50.9%
0.99	6,274	38,565	−41.5%

Table 4: Dedup-threshold sensitivity (near-duplicate removal alone, on pre-curation groups; the group-integrity prune removes a further ≈ 500 to reach the headline 31,865).

Component	<i>n</i>	Human-checked result
Age extraction	100	LLM 89%, regex 71% exact
Group integrity	100	single/multi acc. 0.96
Positive pairs	200	same-person prec. 0.96
Cluster merges	100	merge prec. 0.95
Near-dup removals	100	duplicate prec. 0.93

Table 5: Small manual validation (one annotator; a sanity check, *not* a certified audit; stratified/random/near-threshold sampling). The 100-post counts are illustrative and to be confirmed with the final annotation.

or legal categories), and 45+ is sparse ($n=153$).

At a fixed operating point (per-model global threshold at 1% FMR), +pairs lowers the false-non-match rate in *every* apparent stratum (0–17 0.852 $\rightarrow$ 0.656, apparent-female 0.674 $\rightarrow$ 0.509, apparent-male 0.743 $\rightarrow$ 0.534), with per-stratum logistic calibration slopes uniformly positive (9.0–11.7); the only exception is the small 45+ stratum, whose FMR scatters at the global threshold (sampling noise). The child $\leftrightarrow$ adult gain is corroborated on an independent FG-NET subset (one photo under 13, the other over 25; 195 positives): +pairs raises ROC-AUC 0.684 [0.63, 0.74] $\rightarrow$

Backbone (strength)	Frozen	+pairs	Δ large-gap	Δ our.25+
FaceNet (weak)	0.736	0.848	+0.112	+0.198
ArcFace r50 (weak, diff. arch.)	0.764	0.805	+0.041	+0.133
AdaFace IR-50 (strong)	0.917	0.924	+0.007	+0.004
ArcFace r100 (strong)	0.954	0.915	−0.039	−0.002
AdaFace IR-101 (strong)	0.957	0.931	−0.026	−0.006

Table 6: Headroom: FG-NET large-gap by backbone strength (curated, lr 3e-5). The gain is specific to weak/medium backbones with cross-age headroom.

Train fraction	# identities	FG-NET large-gap	our.25+
frozen	0	0.736	0.640
0.10	$\approx 1,500$	0.846 ± 0.002	0.816 ± 0.003
0.25	$\approx 3,750$	0.852 ± 0.000	0.842 ± 0.002
0.50	$\approx 7,499$	0.856 ± 0.005	0.845 ± 0.002
1.00	$\approx 14,998$	0.848 ± 0.001	0.835 ± 0.003

Table 7: Data-scaling law (FaceNet, curated; val/test fixed; mean $\pm$ std over three seeds).

Model	Age probe $\downarrow$	Identity AUC $\uparrow$
frozen	0.651 ± 0.019	0.836
+pairs	0.629 ± 0.012	0.916
+pairs+disent.	0.632 ± 0.017	0.918

Table 8: Age-leakage linear probe (curated; chance = 0.25). “Age probe” is the balanced accuracy of a linear age classifier read off the identity embedding (lower is better).

0.813 [0.77, 0.86] (+0.129, non-overlapping 95% CIs).

7 Cross-Source and Cross-Platform Transfer

Table 10 reports the VK-trained model evaluated on the independent non-VK Reddit then/now set, and Table 11 the bidirectional transfer (a backbone trained on *only* a small 220-person Reddit source also helps on VK).

8 Calibration and Hard-Negative Mining

Calibration. Raw cosine (Platt) is mis-calibrated across gap bins—over-confident at 15–25 years (ECE 0.052); a $P(\text{same} \mid \text{cosine}, \text{age-gap})$ calibrator fixes every bin (15–25 $\rightarrow$ 0.017) and improves the overall Brier score (0.035 $\rightarrow$ 0.029), yielding a usable, age-aware same-identity probability.

Hard-negative mining. Adding the top-5 most-similar other-person faces per anchor (in a hard cosine band, excluding the near-duplicate/uncertain range) further sharpens large-gap discrimination (Table 12): FG-NET large-gap rises to 0.868 ± 0.004 over three seeds, but easy benchmarks forget *more* and FG-NET overall ROC drops. Hard negatives thus trade general-benchmark accuracy for large-gap separation—useful for a cross-age retrieval setting, not a general verifier.




Figure 2: Apparent-demographic audit: frozen vs. +pairs ROC-AUC across apparent gender/age strata. Every stratum improves; the largest gain is for the youngest (0–17) group. Error bars are 95% paired-bootstrap CIs of the gain.

9 Claims vs. Evidence

Table 13 maps each substantive claim to its evidence and scope.

10 Model and Data Card

Intended use. Research on consent-based cross-age identity matching; the model outputs candidates for human review, never an automatic decision. **Out of scope:** mass identification, surveillance, de-anonymization, any processing of minors’ data without a legal basis. **Training data.** Two public VK “then/now” communities (one platform, Russian-language context); apparent-female-skewed (76.7%); sparse at apparent age 45+; weak same-person supervision from same-post co-occurrence with LLM-parsed caption ages as a secondary cue. **Evaluation data.** LFW, AgeDB-30, CALFW, FG-NET (and its large-gap and child $\leftrightarrow$ adult subsets), an internal curated cross-age test, and an independent non-VK Reddit then/now set. **Metrics.** ROC-AUC (threshold-free; primary), EER, TAR@FAR, and per-stratum calibration. **Ethical considerations.** Sensitive biometric data; no raw faces/posts released; weights/embeddings shared on re-

Stratum	n_{pos}	Frozen	+pairs	Gain	95% CI
overall	5,107	0.836	0.916	+0.080	[+0.075, +0.086]
apparent F	3,790	0.840	0.923	+0.084	[+0.077, +0.090]
apparent M	1,317	0.838	0.897	+0.059	[+0.048, +0.071]
age 0–17	1,123	0.750	0.880	+0.130	[+0.113, +0.146]
age 18–29	3,159	0.872	0.936	+0.064	[+0.058, +0.069]
age 30–44	672	0.822	0.893	+0.070	[+0.054, +0.087]
age 45+	153	0.859	0.893	+0.034	[+0.004, +0.064]

Table 9: Apparent gender/age strata (FaceNet, curated; 95% paired-bootstrap CI of the gain). No stratum degrades.

Reddit (non-VK) test metric	Frozen	+pairs (VK-trained)
Overall ROC-AUC [95% CI]	0.718 [0.70, 0.74]	0.749 [0.73, 0.77]
EER	0.345	0.328
TAR@FAR=1%	0.271	0.279
Test pairs (pos / neg)	1,575 / 1,575	

Table 10: Cross-platform transfer: a VK-trained model on an independent Reddit (non-VK) then/now test set (396 posts $\rightarrow$ 1,575 positive pairs after the multi-person-collage filter).

Train	Test	Frozen	Tuned	Δ
VK	Reddit (overall)	0.718	0.749	+0.031
Reddit	FG-NET large-gap	0.736	0.783	+0.047
Reddit	VK internal 25+	0.640	0.659	+0.019

Table 11: Cross-source transfer, both directions. Reddit $\rightarrow$ VK is a small 220-person pilot supporting source non-specificity, not a headline claim.

Metric	Frozen	+pairs	+pairs+hard-neg
FG-NET large-gap	0.736	0.848	0.868 $\pm$ 0.004
FG-NET ROC	0.896	0.923	0.912 $\pm$ 0.003
LFW accuracy	0.964	0.968	0.925 $\pm$ 0.009
AgeDB-30 ROC	0.953	0.946	0.911 $\pm$ 0.013
CALFW ROC	0.948	0.949	0.898 $\pm$ 0.015

Table 12: Hard-negative mining (FaceNet, curated; external benchmarks). +hard-neg is mean $\pm$ std over three seeds.

quest under a brief research-use agreement; the de-identified dataset available on request (formal documentation if required); minors withheld absent a legal framework; DPIA and datasheet (Geburu et al. 2021) accompany the release. **Caveats.** Targeted large-gap gain only (low-FAR operating points degrade); concentrates in weak/medium backbones; apparent (not self-identified) attributes; audited but not certified supervision.

11 Reproducibility Details

Pipeline. *Usable-face* filtering (detection score, minimum size, Laplacian blur, single target); detection/alignment with InsightFace `buffalo_l` RetinaFace + 5-point ArcFace `norm_crop` (112 px; FaceNet input 160 px); ArcFace `w600k_r50` embeddings for clustering (union-find, cosine $\geq$ 0.85) and dedup (cosine $\geq$ 0.97); an LLM with

versioned, cached prompts for age and group-integrity parsing; a by-person 70/15/15 split (seed 42, balanced negatives, plus an age-matched negative variant). **Training.** Margin-contrastive head fine-tuning (lr 3e-5 weak / 1e-5 strong; early stop on validation AUC; seeds 42/1/2); margin identity-classification objectives (ArcFace/CosFace/SphereFace, and sub-center ArcFace with $K=3$) for the objective comparison; pure-NumPy ROC-AUC/EER/TAR@FAR. **Compute.** The entire study runs on a single laptop: one NVIDIA RTX 3060 Laptop GPU (6 GB), an AMD Ryzen 7 5800H (8C/16T) and 16 GB RAM under Windows 11; Python 3.12, PyTorch 2.12 (CUDA 13.2), InsightFace 1.0 / ONNX Runtime 1.26, scikit-learn 1.9, Matplotlib 3.11. **Availability.** Code, configurations, versioned LLM prompts, frozen splits (by hash), and the de-identified benchmark protocol are released; raw face images and posts are withheld, so data-mining replication is partial (we release cached LLM outputs for deterministic replay). A filled reproducibility checklist accompanies the submission.

Claim	Evidence	Scope / caveat
Mined pairs improve large-gap cross-age verification	FG-NET large-gap 0.736 $\rightarrow$ 0.848, non-overlapping CIs, DeLong $p < 10^{-19}$	weak/medium backbones
Data source > objective choice	contrastive/ArcFace/CosFace coincide (0.848/0.851/0.843); sub-center adds nothing	among objectives tested
Real pairs > synthetic aging	real 0.848 vs. gap-matched FRAN 0.759; native-res control	this protocol
Mechanism is age-shortcut removal	age-matched negatives collapse frozen to ≈ 0.52 ; +pairs recovers to 0.838	internal 25+ test
Source generalizes beyond platform	VK $\rightarrow$ Reddit +0.031; child $\leftrightarrow$ adult FG-NET +0.129	not platform-agnostic
Targeted, not general gain	concentrated on the large-gap subset; easy benchmarks flat (LFW TAR@0.1% 0.814 $\rightarrow$ 0.816)	scope to large-gap retrieval
$\approx$ half of raw positives were near-duplicates	–52% positive count, stable across dedup 0.93–0.97	curation finding

Table 13: Claims-vs-evidence map.

References

- Gebru, T.; Morgenstern, J.; Vecchione, B.; Vaughan, J. W.; Wallach, H.; Daumé III, H.; and Crawford, K. 2021. Datasheets for Datasets. *Commun. ACM*, 64(12): 86–92.
- Huang, Z.; Zhang, J.; and Shan, H. 2021. When Age-Invariant Face Recognition Meets Face Age Synthesis: A Multi-Task Learning Framework. In *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 7282–7291.
- Wang, H.; Gong, D.; Li, Z.; and Liu, W. 2019. Decorrelated Adversarial Learning for Age-Invariant Face Recognition. In *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 3527–3536.
- Wang, H.; Sanchez, V.; and Li, C.-T. 2024. Cross-Age Contrastive Learning for Age-Invariant Face Recognition. In *Proc. IEEE Int. Conf. Acoust. Speech Signal Process. (ICASSP)*. ArXiv:2312.11195.
- Wang, H.; Sanchez, V.; and Li, C.-T. 2025. From Age Estimation to Age-Invariant Face Recognition: Generalized Age Feature Extraction Using Order-Enhanced Contrastive Learning. *arXiv preprint arXiv:2501.01760*.
- Wang, Y.; Gong, D.; Zhou, Z.; Ji, X.; Wang, H.; Li, Z.; Liu, W.; and Mei, T. 2018. Orthogonal Deep Features Decomposition for Age-Invariant Face Recognition. In *Proc. Eur. Conf. Comput. Vis. (ECCV)*, 738–753.
- Yao, W.; Farooq, M. A.; Lemley, J.; and Corcoran, P. 2024. Synthetic Face Ageing: Evaluation, Analysis and Facilitation of Age-Robust Facial Recognition Algorithms. *arXiv preprint arXiv:2406.06932*.